

Federated Single-Index Bandits

Cooperative learning when the reward function itself is unknown

Presented by:

Name 1 (Roll No.) Dept. of CSE

Name 2 (Roll No.) Dept. of CSE

Course: CS6007: Multi-Agent Machine Learning

Instructor: Prof. Avishek Ghosh

IIT Bombay

28 September 2026

The assumption everybody makes, and why it breaks

Contextual bandits assume we know how context maps to reward:

$$\mathbb{E}[y \mid x] = \underbrace{\langle x, \theta_* \rangle}_{\text{linear bandit}} \quad \text{or} \quad \underbrace{f(\langle x, \theta_* \rangle)}_{\substack{\text{generalized linear} \\ f \text{ known}}}$$

¹Ghosh, Chowdhury & Gopalan, *Misspecified Linear Bandits*, AAAI 2017.

The assumption everybody makes, and why it breaks

Contextual bandits assume we know how context maps to reward:

$$\mathbb{E}[y \mid x] = \underbrace{\langle x, \theta_* \rangle}_{\text{linear bandit}} \quad \text{or} \quad \underbrace{f(\langle x, \theta_* \rangle)}_{\substack{\text{generalized linear} \\ f \text{ known}}}$$

But the link function is essentially never known

Dose-response, congestion, click-through and user-satisfaction curves are non-monotone and unknown in advance.

¹Ghosh, Chowdhury & Gopalan, *Misspecified Linear Bandits*, AAAI 2017.

The assumption everybody makes, and why it breaks

Contextual bandits assume we know how context maps to reward:

$$\mathbb{E}[y \mid x] = \underbrace{\langle x, \theta_* \rangle}_{\text{linear bandit}} \quad \text{or} \quad \underbrace{f(\langle x, \theta_* \rangle)}_{\substack{\text{generalized linear} \\ f \text{ known}}}$$

But the link function is essentially never known

Dose-response, congestion, click-through and user-satisfaction curves are non-monotone and unknown in advance.

This is not a mild modelling error

Misspecifying f makes state-of-the-art algorithms incur **linear regret** — they do not degrade gracefully, they fail.¹

¹Ghosh, Chowdhury & Gopalan, *Misspecified Linear Bandits*, AAAI 2017.

Problem statement: the single-index bandit (SIB)

Model

At round t the learner sees K arms $\mathcal{X}_t = \{x_{t,a}\}_{a \in [K]}$ drawn i.i.d. from a density p on $\mathbb{R}^d$; it pulls x_t and observes

$$y_t = f(\langle x_t, \theta_* \rangle) + \eta_t, \quad \eta_t \text{ } \sigma\text{-sub-Gaussian,}$$

with **both** $\theta_* \in \mathbb{R}^d$ **and** $f : \mathbb{R} \rightarrow \mathbb{R}$ **unknown**. Normalise $\|\theta_*\|_1 = 1$.

Problem statement: the single-index bandit (SIB)

Model

At round t the learner sees K arms $\mathcal{X}_t = \{x_{t,a}\}_{a \in [K]}$ drawn i.i.d. from a density p on $\mathbb{R}^d$; it pulls x_t and observes

$$y_t = f(\langle x_t, \theta_* \rangle) + \eta_t, \quad \eta_t \text{ } \sigma\text{-sub-Gaussian,}$$

with **both** $\theta_* \in \mathbb{R}^d$ **and** $f : \mathbb{R} \rightarrow \mathbb{R}$ **unknown**. Normalise $\|\theta_*\|_1 = 1$.

f need not be monotone $\Rightarrow$ the best arm is *not* the one with the largest projection.

Problem statement: the single-index bandit (SIB)

Model

At round t the learner sees K arms $\mathcal{X}_t = \{x_{t,a}\}_{a \in [K]}$ drawn i.i.d. from a density p on $\mathbb{R}^d$; it pulls x_t and observes

$$y_t = f(\langle x_t, \theta_* \rangle) + \eta_t, \quad \eta_t \text{ } \sigma\text{-sub-Gaussian,}$$

with **both** $\theta_* \in \mathbb{R}^d$ **and** $f : \mathbb{R} \rightarrow \mathbb{R}$ **unknown**. Normalise $\|\theta_*\|_1 = 1$.

f need not be monotone $\Rightarrow$ the best arm is *not* the one with the largest projection.

Key tool — Stein's identity, with score $S(x) = -\nabla \log p(x)$

$$\mathbb{E}[y S(x)] = \mathbb{E}[f'(\langle x, \theta_* \rangle)] \theta_* = \mu_* \theta_*$$

The direction is recoverable by a **plain sample average** — with no knowledge of f whatsoever.

Where the literature stands

	unknown link f	multi- agent	safety constr.	optimal rate
Abbasi-Yadkori et al. '11 (OFUL)	✗	✗	✗	—
Chowdhury & Gopalan '17 (IGP-UCB)	✓	✗	✗	—
Dubey & Pentland '20 (Coop-KernelUCB)	✓	✓	✗	—
Amani & Thrampoulidis '21 (Safe-DLUCB)	✗	✓	✓	—
Kang et al. '26 (ESTOR / GSTOR)	✓	✗	✗	✗ $T^{3/4}$
Arya et al. '26 (kernel SIB)	✓	✗	✗	—
Dey, Bhore & Ghosh '26 (ZoomSIB-UCB)	✓	✗	✗	✓ $\tilde{\Theta}(T^{2/3})$
This project (Fed-ZoomSIB)	✓	✓	(ext.)	?

Where the literature stands

	unknown link f	multi- agent	safety constr.	optimal rate
Abbasi-Yadkori et al. '11 (OFUL)	✗	✗	✗	—
Chowdhury & Gopalan '17 (IGP-UCB)	✓	✗	✗	—
Dubey & Pentland '20 (Coop-KernelUCB)	✓	✓	✗	—
Amani & Thrampoulidis '21 (Safe-DLUCB)	✗	✓	✓	—
Kang et al. '26 (ESTOR / GSTOR)	✓	✗	✗	✗ $T^{3/4}$
Arya et al. '26 (kernel SIB)	✓	✗	✗	—
Dey, Bhore & Ghosh '26 (ZoomSIB-UCB)	✓	✗	✗	✓ $\tilde{\Theta}(T^{2/3})$
This project (Fed-ZoomSIB)	✓	✓	(ext.)	?

Every column has been solved individually.

No work occupies the last row: single-index bandits have never been federated.

Our problem: what does collaboration buy?

Setting

N agents on a connected graph G with communication matrix $\mathbf{P}$. Each round, every agent i sees its own arm set, pulls $x_{i,t}$, observes $y_{i,t} = f(\langle x_{i,t}, \theta_* \rangle) + \eta_{i,t}$. f and θ_* are **shared and unknown**. Agents exchange $\leq B$ values with neighbours every C rounds.

Objective — network regret

$$R_T(N) = \sum_{t=1}^T \sum_{i=1}^N \left[f(\langle x_{i,t}^*, \theta_* \rangle) - f(\langle x_{i,t}, \theta_* \rangle) \right]$$

Our problem: what does collaboration buy?

Setting

N agents on a connected graph G with communication matrix $\mathbf{P}$. Each round, every agent i sees its own arm set, pulls $x_{i,t}$, observes $y_{i,t} = f(\langle x_{i,t}, \theta_* \rangle) + \eta_{i,t}$. f and θ_* are **shared and unknown**. Agents exchange $\leq B$ values with neighbours every C rounds.

Objective — network regret

$$R_T(N) = \sum_{t=1}^T \sum_{i=1}^N \left[f(\langle x_{i,t}^*, \theta_* \rangle) - f(\langle x_{i,t}, \theta_* \rangle) \right]$$

How much does cooperation help when the reward function itself must be learned — and at what communication cost?

Why this algorithm invites federation

ZOOMSIB-UCB (Dey, Bhore & Ghosh '26) is two-phase, and each phase is shareable — but by a *different mechanism*.

Phase 1 — estimate θ_*

Truncated Stein estimator over T_0 rounds.

It is a **sample mean** — so averaging local estimators is **algebraically identical** to the centralized estimator on pooled data.

Exact federation, **one** message of d numbers, no consensus iterations, no approximation.

Phase 2 — UCB over bins

$N_{\text{bins}} = O(T^{1/3})$ bins on the index line; *sleeping* UCB.

The entire learner state is **two scalars per bin** (n_j, S_j) — so sharing it is a standard cooperative MAB over the graph G .

Cheap to communicate, but the pooling is statistical, not exact.

Why this algorithm invites federation

ZOOMSIB-UCB (Dey, Bhore & Ghosh '26) is two-phase, and each phase is shareable — but by a *different mechanism*.

Phase 1 — estimate θ_*

Truncated Stein estimator over T_0 rounds.

It is a **sample mean** — so averaging local estimators is **algebraically identical** to the centralized estimator on pooled data.

Exact federation, **one** message of d numbers, no consensus iterations, no approximation.

Phase 2 — UCB over bins

$N_{\text{bins}} = O(T^{1/3})$ bins on the index line; *sleeping* UCB.

The entire learner state is **two scalars per bin** (n_j, S_j) — so sharing it is a standard cooperative MAB over the graph G .

Cheap to communicate, but the pooling is statistical, not exact.

Prop. 4.7 of Dey et al. is a **modular reduction**: *any* sleeping-bandit subroutine may be substituted into Phase 2 — and a cooperative one is exactly that.

What we set out to answer

We are *not* claiming a rate in advance. The algorithm splits into two independently shareable pieces; how much each one buys is exactly what we intend to measure.

Research questions

- 1 Does sharing reduce **per-agent** regret in a single-index bandit, and by how much?
- 2 The two phases pool very differently — an *exact* average versus a statistical one. **Does the benefit split unevenly between them?**
- 3 **How little communication is enough?** Phase 1 needs a single message; can Phase 2 be made nearly as cheap?
- 4 How does any gain scale with N , with dimension d , and with the graph topology (spectral gap $1 - |\lambda_2|$)?

What we set out to answer

We are *not* claiming a rate in advance. The algorithm splits into two independently shareable pieces; how much each one buys is exactly what we intend to measure.

Research questions

- 1 Does sharing reduce **per-agent** regret in a single-index bandit, and by how much?
- 2 The two phases pool very differently — an *exact* average versus a statistical one. **Does the benefit split unevenly between them?**
- 3 **How little communication is enough?** Phase 1 needs a single message; can Phase 2 be made nearly as cheap?
- 4 How does any gain scale with N , with dimension d , and with the graph topology (spectral gap $1 - |\lambda_2|$)?

If time permits. (a) Safety: an unknown constraint $h(\langle x, \mu_* \rangle) \leq c$. **(b) Clustered indices:** agents in M unknown groups, IFCA-style.

Project plan

Window	Work	Deliverable
29 Sep – 10 Oct	Reproduce single-agent ZOOMSIB-UCB; baselines: Random, LinUCB, ESTOR, GSTOR, IGP-UCB	Regret curve matching Dey et al. Fig. 1
11 – 22 Oct	Fed-ZoomSIB : one-shot Stein averaging + cooperative sleeping-UCB over graph topologies	Speed-up plots; federated concentration lemma
23 – 29 Oct	Communication ablations: interval C , event-triggered, quantization	Regret per transmitted byte
30 Oct – 5 Nov	One extension: safety <i>or</i> clustered indices	Violation-rate / negative-transfer study
5 – 12 Nov	Final deck + report, ≥ 20 seeds with CIs	Submission

Project plan

Window	Work	Deliverable
29 Sep – 10 Oct	Reproduce single-agent ZOOMSIB-UCB; baselines: Random, LinUCB, ESTOR, GSTOR, IGP-UCB	Regret curve matching Dey et al. Fig. 1
11 – 22 Oct	Fed-ZoomSIB : one-shot Stein averaging + cooperative sleeping-UCB over graph topologies	Speed-up plots; federated concentration lemma
23 – 29 Oct	Communication ablations: interval C , event-triggered, quantization	Regret per transmitted byte
30 Oct – 5 Nov	One extension: safety <i>or</i> clustered indices	Violation-rate / negative-transfer study
5 – 12 Nov	Final deck + report, ≥ 20 seeds with CIs	Submission

Evaluation: network & per-agent regret; samples to fixed accuracy; messages and bytes; sensitivity to spectral gap $1 - |\lambda_2|$; sweeps over N, d, K , topology, link shape. Synthetic + KDD Cup 99 and Forest Cover Type, as in Dey et al.

References

- ① D. Dey, S. Bhore, A. Ghosh. *Optimal Regret for Single Index Bandits*. arXiv:2605.09454, 2026.
- ② Y. Kang, M. Liu, B. Yi, J. Lyu, Z. Zhang, D. Zhou, Y. Li. *Single Index Bandits: Generalized Linear Contextual Bandits with Unknown Reward Functions*. ICLR, 2026.
- ③ A. Dubey, A. Pentland. *Kernel Methods for Cooperative Multi-Agent Contextual Bandits*. ICML, 2020.
- ④ S. Amani, C. Thrampoulidis. *Decentralized Multi-Agent Linear Bandits with Safety Constraints*. AAAI, 35(8):6627–6635, 2021.
- ⑤ S. Arya, S. Bhattacharjee, B. K. Sriperumbudur. *Kernel Single-Index Bandits: Estimation, Inference, and Learning*. arXiv:2603.18938, 2026.
- ⑥ S. R. Chowdhury, A. Gopalan. *On Kernelized Multi-armed Bandits*. ICML, 2017.
- ⑦ A. Ghosh, J. Chung, D. Yin, K. Ramchandran. *An Efficient Framework for Clustered Federated Learning*. NeurIPS, 2020.
- ⑧ A. Ghosh, S. R. Chowdhury, A. Gopalan. *Misspecified Linear Bandits*. AAAI, 2017.

Thank You!